

Report: Project Data Exploration and Analysis

-
Charles Meyer

In this project, we leverage the skills acquired in class to conduct data analysis using diverse technologies. Our objective is to explore a dataset of our choosing, define a compelling business objective, and employ a series of analytical techniques to achieve meaningful insights. This report outlines our approach, challenges encountered, and potential enhancements for future iterations.

This project aims to develop a dashboard for runners, providing various metrics, graphs, and statistics to help them manage their training and track performance over time. By leveraging big data technologies, we process running-related datasets to extract insights and present them in an interactive and user-friendly format. The project consists of two main components: a Scala-based data processing pipeline using Apache Spark for efficient data handling and a Python-based dashboard built with Dash for visualization and interaction.

The initial stages of the project involved preliminary data exploration and basic data cleaning using a Jupyter notebook. During this phase, we identified missing values in the country column, opting not to remove these rows due to completeness in other columns.




Figure 1 : Architecture overview

The Scala component of our project is structured around several key classes: **SparkSessionBuilder** for initiating Spark sessions, **DataReader** for CSV file ingestion, **DataWriter** for output management, **DataProcessor** for core data manipulation functions, and **Main** for orchestrating these processes. We chose to forego Hadoop Distributed File System (HDFS) as our dataset sizes did not necessitate its complexity.

For data visualization and interaction, we employed the **Dash** library in Python, leveraging its familiarity and efficiency for creating our dashboard. The project was organized into multiple files including **run.py** for application startup, **graph.py** for graph creation before integrating them into the dashboard, **app/callback.py** for callback functions, **app/app.py** for page architecture, **app/dataHandler.py** for data management, and **app/assets/** for static resources.

While no specific results are presented due to the project's nature, certain dataset limitations were identified. These include the absence of detailed run-specific metrics, such as kilometer-by-kilometer pace analysis or comprehensive heart rate data, which hindered

deeper analyses. Potential improvements identified include transitioning to a relational database like PostgreSQL or SQLite to reduce dashboard latency and enable backup systems. Furthermore, automating data processing upon new file reception is recommended to streamline operations.

Looking forward, integrating a machine learning model to forecast daily injury risks for runners stands as a promising future enhancement.